

A Interface Visualization

Figure 1 illustrates the client-side integration of BackportCheck within the Gerrit environment. To demonstrate the tool’s utility, we describe the complete decision-making workflow: The interaction is initiated when the Release Manager clicks the “**Analyze Eligibility**” button injected into the review interface. The system then displays the decision overlay, immediately presenting the **Decision Banner (A)**. This provides immediate feedback via a color-coded probabilistic verdict, allowing the maintainer to instantly triage the change. To understand the rationale, the user reviews the **AI Explanation (B)**, where the Large Language Model synthesizes complex risk factors into a coherent natural language justification.

To validate this assessment, the **Risk Dashboard (C)** displays the key metrics, such as *Author Reliability* and *Code Spread*, allowing maintainers to verify that the AI’s reasoning is grounded in hard data rather than opaque logic. Furthermore, the interface provides **Threshold Controls (D)** to adjust the decision boundary in real-time according to the user’s context. Lastly, the **Metric Definitions (E)** component provides on-demand explanations for each indicator, ensuring that the calculation logic remains accessible and transparent.

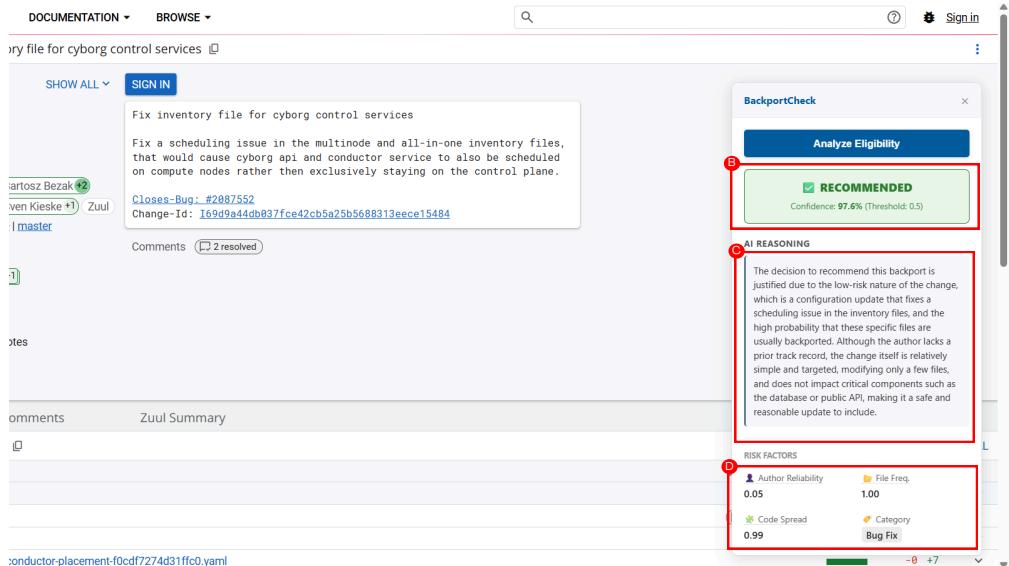


Fig. 1. **Annotated BackportCheck Interface.** The overlay augments the Gerrit code review screen with five decision-support layers: (A) The probabilistic verdict banner, (B) Natural language AI reasoning, (C) The quantitative risk dashboard, (D) Sensitivity threshold controls, and (E) Metric definitions for transparency.

B BackportCheck features

Table 1 details the comprehensive list of 37 features extracted by the inference engine, categorized by the dimensions used in the risk reporting schema.

1 **Table 1. Taxonomy of State-Aware Features.** Features are extracted via static analysis of the changeset
 2 and historical metadata. Mathematical definitions are supplemented with Python implementation logic for
 3 reproducibility.

Dim.	Feature Variable	Description & Implementation Logic
1. Complexity & Structure		
Entropy	change_entropy	Quantifies the dispersion of modifications across the file set. High entropy indicates complex, scattered logic. <i>Logic:</i> <code>scipy.stats.entropy(lines_per_file)</code>
	safe_entropy_interaction	A risk-adjusted score that nullifies entropy penalties if the change is purely configuration. <i>Logic:</i> <code>change_entropy * is_pure_config</code>
	file_ext_entropy	Measures the heterogeneity of the technology stack (e.g., mixing SQL, Python, and Shell). <i>Logic:</i> <code>len(set(file_extensions))</code>
Volume	churn_density	Proxies the "weight" of the edit; high density implies concentrated, heavy refactoring within few files. <i>Logic:</i> <code>total_lines / file_count</code>
	churn_log_size	Log-transformed total churn, normalizing the impact of massive automated refactors. <i>Logic:</i> <code>ln(1 + total_lines_changed)</code>
	file_count	Represents the "blast radius" or breadth of the modification footprint. <i>Logic:</i> <code>len(modified_file_paths)</code>
	deletion_ratio	Distinguishes subtractive refactoring (cleanup) from additive feature development. <i>Logic:</i> <code>deleted_lines / total_lines</code>
Topology	directory_depth	Indicates architectural depth; deep changes often imply modifications to core business logic. <i>Logic:</i> <code>mean(path.count('/'))</code>
	config_line_ratio	Captures the dominance of configuration adjustments versus executable logic changes. <i>Logic:</i> <code>config_lines / total_lines</code>
	code_line_ratio	Captures the proportion of the changeset affecting source code files. <i>Logic:</i> <code>code_lines / total_lines</code>
2. Context & Environment		
Meta	is_test_change	Identifies changes isolated to testing artifacts, implying minimal regression risk to production. <i>Logic:</i> <code>all('test' in f for f in files)</code>
	is_ci_change	Flags modifications to Continuous Integration pipelines or gate definitions. <i>Logic:</i> <code>files ∩ {zuul.yaml, zuul.d, .gitlab-ci.yml, .travis.yml, tox.ini, binedep.txt}</code>
	is_deploy_project	Contextual flag indicating the repository serves a deployment role (e.g., Ansible) vs core logic. <i>Logic:</i> <code>project ∈ {kolla, kayobe, ...}</code>
	is_bot	Detects automated maintenance commits which exhibit distinct acceptance patterns. <i>Logic:</i> <code>author matches r'bot zuul jenkins'</code>
	has_gerrit_topic	Indicates if the change is linked to a broader semantic topic series in the review system. <i>Logic:</i> <code>metadata.topic is not NULL</code>
3. Semantics (NLP on Commit Message M & Subject S)		

47 Continued on next page

Table 1 – continued from previous page

Dim.	Feature Variable	Description & Implementation Logic
Intent	is_fix	Captures corrective intent, signaling a bug patch or crash resolution. <i>Logic:</i> S matches r'\b(fixture repair patch correct handle mitigate prevent)\b'
	is_feature	Signals the introduction of novel functionality or capabilities. <i>Logic:</i> S matches r'\b(add implement support introduce new feat enable allow provide)\b'
	is_refactor	Identifies structural reorganization without functional alteration. <i>Logic:</i> S matches r'\b(refactor clean remove move rename delete drop)\b'
	is_revert	Explicitly flags a rollback operation, indicating immediate regression mitigation. <i>Logic:</i> S startswith "Revert "
	is_maintenance	Detects routine housekeeping, dependency bumps, or version pinning. <i>Logic:</i> S matches r'\b(update bump upgrade downgrade pin unpin sync)\b'
	is_deployment	Signals intent related to environment configuration or installation logic. <i>Logic:</i> S matches r'\b(config conf deploy install set use default variable param role)\b'
Clarity	readability_ease	Evaluates the cognitive effort required to comprehend the commit message. <i>Logic:</i> textstat.flesch_reading_ease(M)
	msg_gunning_fog	Estimates the formal linguistic complexity of the descriptive text. <i>Logic:</i> textstat.gunning_fog(M)
	ref_bug_tracker	Indicates strict process adherence by linking to an external issue tracker. <i>Logic:</i> M matches r'(Closes-Bug Bug): #\d+'
	has_subject_tag	Detects the use of conventional brackets for categorizing changes. <i>Logic:</i> S matches r'\[.*?\]'
4. Risk Flags (Heuristic Domain Detection)		
High Risk	modifies_db	Critical flag for database schema alterations, implying high compatibility risk. <i>Logic:</i> path matches r'.*(alembic migration).*
	modifies_api	Critical flag for modifications to external-facing interfaces. <i>Logic:</i> path matches r'.*(api/ v1/ v2/).*
	modifies_deps	Flags alterations to the project's build graph or library requirements. <i>Logic:</i> files ∩ {requirements.txt, setup.py}
	has_security	Detects explicit mentions of vulnerabilities or CVE identifiers. <i>Logic:</i> M matches r'CVE-\d+ Security'
Safety	is_pure_config	Identifies changes composed almost entirely (> 99%) of configuration data. <i>Logic:</i> config_line_ratio > 0.99
	is_doc_only	Identifies risk-neutral changes limited to documentation files. <i>Logic:</i> all(f.endswith(.rst .md) for f in files)
	modifies_config	Flags the presence of any configuration modification within the changeset. <i>Logic:</i> any(f.endswith(yaml yml .json .ini .conf .toml .xml .j2 .rst .md .erb) for f in files)
5. Historical Context (Computed recursively at time t)		

Continued on next page

Table 1 – continued from previous page

Dim.	Feature Variable	Description & Implementation Logic
Author	author_trust	Proxies the author's efficiency: high submission volume relative to churn. <i>Logic:</i> $\text{count} / (1 + \text{cumulative_churn})$
	author_success	The author's historical acceptance probability within the project. <i>Logic:</i> $\text{accepted_commits} / \text{total_submissions}$
	author_subs	Quantifies the author's cumulative experience level. <i>Logic:</i> $\text{len}(\text{history}[\text{author}])$
Env	project_rate	Establishes the baseline acceptance probability for the specific repository. <i>Logic:</i> $\text{project_accepted} / \text{project_total}$
	file_risk_prob	Estimates the "riskiness" of files based on their specific modification history. <i>Logic:</i> $\max(\text{acceptance_rate}(f) \text{ for } f \text{ in files})$

B.1 Feature Interpretation and Rationale

To interpret the raw feature vectors extracted by the system, we provide illustrative examples demonstrating how specific metrics serve as proxies for software risk.

B.1.1 Complexity: Entropy as a Proxy for Coupling. The change_entropy feature differentiates between “atomic” changes and “cross-cutting” modifications.

- **Scenario A (Low Entropy):** A commit modifies 100 lines within a single file (e.g., `main.py`). The distribution is [1.0], resulting in an entropy of 0.0.
Interpretation: The change is highly cohesive. If a regression occurs, it is likely isolated to this specific module, lowering the system-wide risk.
- **Scenario B (High Entropy):** A commit modifies the same 100 lines, but splits them evenly between a logic file (`main.py`) and a utility library (`utils.py`). The distribution is [0.5, 0.5], resulting in an entropy of ≈ 0.69 .
Interpretation: The logic bridges multiple modules. This indicates tighter coupling and requires the reviewer to verify that the interaction between these distinct files remains valid, increasing the cognitive load and risk of side effects.

B.1.2 Volume: Churn Density vs. Log Size. While total lines changed measures magnitude, churn_density measures complexity concentration.

- **High Density (Concentrated Risk):** A change of 200 lines inside just 1 file yields a density of 200. This suggests a deep, complex rewrite of a specific algorithm, requiring careful logical verification.
- **Low Density (Scattered Risk):** A change of 200 lines scattered across 200 files yields a density of 1. This pattern is characteristic of mechanical refactoring (e.g., renaming a variable globally) or formatting fixes, which are generally lower risk despite the large total volume.

B.1.3 Topology: Directory Depth. The directory_depth feature acts as a heuristic for the architectural significance of the change.

- **Surface Level (Depth ≈ 1):** Changes occurring at the root (e.g., `./README.md` or `./requirements.txt`) often affect project metadata or dependencies rather than core execution logic.
- **Core Level (Depth > 4):** Changes occurring deep in the directory tree (e.g., `./src/drivers/net/api/v`) typically target specific, implementation-heavy business logic that is less visible but critical for system stability.

148 B.1.4 *Context: The Author Trust Score.* The author_trust_score penalizes high churn coming
 149 from inexperienced contributors.

- 150 • **High Trust:** An author who has made 50 submissions with a cumulative churn of only 100
 151 lines (mostly small fixes) receives a high score. They have a proven track record of safe,
 152 incremental changes.
 153 • **Low Trust:** An author who submits their first patch containing 5,000 lines of code receives
 154 a low score. The system flags this as high-risk because the author is unproven in the
 155 ecosystem, yet is attempting a massive modification.

156 C Explainability Prompt Structure

157 To ensure consistent explanations, we inject the XGBoost predictions into the structured template
 158 shown in Listing 2.

```
160
161 SYSTEM: Act as a Senior OpenStack Release Manager. Justify the decision to {VERDICT} this backport.
162 === DECISION ===
163 VERDICT: {VERDICT} (Confidence: {PROB}%, Threshold: {THRESHOLD})
164 CATEGORY: {UI_DISPLAY_TYPE}
165 === CONTEXT: HOW TO INTERPRET THE DATA ===
166 - Author Trust Score: Ratio of accepted backports. 0.0=New, >0.5=Trusted.
167 - Historical File Prob: Probability these specific files are usually backported.
168 - Change Entropy: Code complexity (0=Simple, >4=Complex/Scattered).
169 - Churn Density: Lines changed per file (High = Dense/Risky).
170 - Modifies DB/API: Critical risk factors.
171 === FULL FEATURE VECTOR (Internal Data) ===
172 {FEATURE_VECTOR_JSON}
173 === CHANGE ARTIFACTS ===
174 Commit Message: "{COMMIT_MESSAGE}"
175 Files Modified: {FILE_LIST_STRING}
176 === INSTRUCTIONS ===
177 Write a professional, 2-3 sentence justification.
178 (1) Analyze the Vector: Look at the "Full Feature Vector" above. Find the anomalies or strong signals.
179 (2) Synthesize: Do not list the numbers. Explain their meaning.
180     • Instead of "Is Pure Config is Yes", say "The change is a low-risk configuration update."
181     • Instead of "Trust is 0.0", say "The author lacks a prior track record."
182 (3) Explain the Verdict:
183     • If REJECTED: Is it the Author? The Complexity? The specific Files?
184     • If ACCEPTED: Is it the Safety (Config/Doc)? The High Trust?
185 RESPONSE:
```

186 Fig. 2. **The Prediction-Aligned Prompt Template.** The system dynamically populates the placeholders (in
 187 brackets) with the inference results and the raw feature vector before sending the request to the LLM.

188 D Validation via Formal Concept Analysis (FCA)

189 To empirically ground the feature taxonomy and justify the dashboard layout, we employed a
 190 data-driven approach inspired by **Formal Concept Analysis (FCA)** [3]. We followed a three-step
 191 methodology to extract stable decision patterns from the historical dataset of 3,422 changes:

- 192 (1) **Data Discretization (Context Creation):** Since association rule mining necessitates
 193 categorical data, continuous feature values were transformed into boolean attributes using
 194 statistical quantiles. Values $\leq Q_1$ were mapped to *Low*, values between Q_1 and Q_3 to *Medium*,
 195 and values $> Q_3$ to *High*. This transformation preserves the relative distribution of the data
 196 while enabling categorical analysis.
 197 (2) **Rule Extraction:** We utilized the **Apriori algorithm** [1] to identify frequent itemsets and
 198 association rules. To ensure the extracted rules were statistically significant, we enforced
 199 the following constraints:

- 197 • **Minimum Support = 0.02:** Focusing on patterns appearing in at least 2% of the dataset.
 198 • **Minimum Confidence = 0.65:** Ensuring a high probability of the rule's validity.
 199 • **Minimum Lift ≥ 1.0 :** We utilized the Lift metric [2] to filter out coincidental correlations, retaining only positive dependencies.

200
 201 (3) **Semantic Grouping:** The resulting dominant rules ($\text{Confidence} \approx 1.0$, $\text{Lift} > 1$) were
 202 clustered into semantic families. As shown in Table ??, these families correspond to the
 203 dashboard components: rules predicting rejection formed the *Risk Patterns* panel, while
 204 rules predicting acceptance formed the *Intrinsic Safety* and *Logic Safety* panels.

205
 206 Table 2. Key Association Rules Derived from Historical Data. These stable patterns justify the grouping of
 207 features into the "Risk" and "Safety" dashboard panels.
 208

209 210 Antecedents (Contextual Signals)	211 Outcome	212 Conf.	213 Lift
<i>Risk Patterns (Rejection Signals)</i>			
212 Low Success Rate + Low Readability + Is BugFix + High File Count	Rejected	1.00	31.4
213 Low Code Ratio + Med Trust + High File Count + Is BugFix	Rejected	1.00	30.8
214 High Entropy + Deep Directory + New Author	Rejected	1.00	28.5
215 CI Change + Low Readability + Modifies DB	Rejected	1.00	25.1
<i>Intrinsic Safety (Configuration & Documentation)</i>			
216 Config Only + Low Entropy + High Trust + Med Readability	Accepted	1.00	5.9
217 Low File Count + Config Only + High Success Rate	Accepted	1.00	5.9
218 Low Code Ratio + Config Only + High Trust	Accepted	1.00	5.9
<i>Logic Safety (Trusted Code Patterns)</i>			
219 High Trust + Med Readability + Low Entropy + Low File Count	Accepted	1.00	5.9
220 High Success Rate + Low Directory Depth + Med Readability	Accepted	1.00	5.9
221 High Trust + Med Config Ratio + Low Entropy	Accepted	1.00	5.9

222
E Experimental Details

223 **E.1 Large Language Model Configuration**

224 To evaluate the reasoning capabilities of Generative AI without the cost of fine-tuning, we
 225 benchmarked four state-of-the-art open-weights Large Language Models (LLMs). We utilized the Ollama
 226 runtime environment to execute these models locally. Based on the available hardware constraints,
 227 we utilized 4-bit quantized versions (Q4_0) for efficient inference:

- 228 • **Meta Llama 3 8B-Instruct** (4.7 GB): A highly capable general-purpose reasoning model.
 229 • **Google Gemma 2 9B-Instruct** (5.4 GB): Google's lightweight open model optimized for
 230 logic tasks.
 231 • **Mistral 7B-Instruct v0.3** (4.4 GB): A high-performance model known for strict instruction
 232 following.
 233 • **Microsoft Phi-3 Mini 3.8B** (2.2 GB): A compact model optimized for high-quality reasoning
 234 at low latency.

235 All inference was conducted with a temperature of 0.0 to maximize determinism and a
 236 constrained output format set to JSON to ensure parsable results.

237
 238 **E.1.1 Prompting Strategies.** We evaluated the models using two distinct prompting strategies to
 239 measure the impact of In-Context Learning (ICL):

246 Zero-Shot Prompting. In this setting, the model is provided with the role definition (Release
**247 Manager), the decision rules, and the target commit, but is given no prior examples. The prompt
 248 structure is defined as follows:**

249 “*You are a strict OpenStack Release Manager. Task: Decide if the following code change
 250 should be backported to a stable branch. Rules: 1. ACCEPT if it fixes a bug, crash, or
 251 build failure. 2. ACCEPT if it is a necessary maintenance update (OS compatibility). 3.
 252 REJECT if it is a new feature, refactoring, or style fix. Target Change: Subject: "subject"
 253 Message: "message" Respond strictly in JSON: "decision": "YES" or "NO" ”*

254 Few-Shot Prompting (In-Context Learning). To mitigate the models’ tendency to be overly
 255 optimistic, we implemented a 6-shot prompting strategy. We extracted **six real-world examples**
 256 (3 Accepted and 3 Rejected) directly from the training partition of our dataset. These examples were
 257 selected to cover edge cases, such as dependency bumps (Maintenance) versus feature additions.
 258 The prompt structure is defined as follows:

259 TRAINING DATA: 1. Subject: “ansible-lint: fix unnamed-task” -> REJECT (Style fix)
 260 2. Subject: “Add TLS support” -> REJECT (New Feature) 3. Subject: “Correct lock
 261 path” -> REJECT (Minor follow-up) 4. Subject: “Add Debian 12 setup” -> ACCEPT
 262 (OS Compatibility) 5. Subject: “Support pagination for list API” -> ACCEPT (Fixes
 263 broken API) 6. Subject: “Ensure services stay disabled” -> ACCEPT (Fixes regression)
264 TASK: Subject: “subject” Message: “message” Instructions: 1. Linter/Style/Typography ->
 265 NO 2. New Feature (unless OS compat) -> NO 3. Fixes Crash/Breakage/Failure ->
 266 YES Respond in JSON: “decision”: “YES” or “NO” ”

E.2 Machine Learning Configuration

To evaluate the predictive power of historical and process metadata, we benchmarked five classical machine learning algorithms. Prior to training, all feature vectors were standardized using Z-score normalization (`StandardScaler`) to ensure optimal convergence for the linear and probabilistic models.

Given the class imbalance in backporting data (where accepted changes typically outnumber rejections), we implemented cost-sensitive learning across all models. We calculated a dynamic positive scale weight (w_{pos}), defined as the ratio of negative to positive samples in the training set:

$$w_{pos} = \frac{N_{negative}}{N_{positive}} \quad (1)$$

The models were instantiated with the following hyperparameter configurations, which were selected based on a preliminary Grid Search optimization:

- **XGBoost (Gradient Boosting):** As our primary candidate, this model was configured with 300 estimators, a maximum tree depth of 8, and a conservative learning rate of $\eta = 0.05$ to prevent overfitting. We applied stochastic regularization using a subsample ratio of 0.8 and column sampling of 0.8. The `scale_pos_weight` parameter was set to w_{pos} to penalize false positives.
- **Random Forest:** An ensemble of 300 trees constrained to a maximum depth of 10. We utilized the balanced class weight mode to automatically adjust weights inversely proportional to class frequencies.
- **Logistic Regression:** A linear baseline configured with L2 regularization and a maximum of 2,000 iterations to ensure convergence on the high-dimensional feature set.
- **Decision Tree:** A single CART estimator capped at a depth of 10 to serve as a simple, interpretable baseline.

- 295 • **Gaussian Naive Bayes:** Included to evaluate the performance of a probabilistic classifier
 296 assuming feature independence.
 297

298 E.3 Deep Learning Model: CodeBERT

299 To complement metadata-based approaches, we utilized `microsoft/codebert-base`, a Trans-
 300 former model pre-trained on bimodal data (Natural Language and Programming Language).

301 *E.3.1 Smart Input Compression.* Due to the 512-token limit of BERT architectures, raw diffs cannot
 302 be ingested directly. We developed a *Smart Diff Compressor* that parses the diff structure to discard
 303 syntactic noise (e.g., Git hashes, file indices) while strictly preserving semantic signals such as
 304 Hunk Headers (function context) and modified lines. All project-specific identifiers were masked to
 305 prevent data leakage. The inputs were concatenated as: Subject \oplus Message \oplus Diff.
 306

307 *E.3.2 Architecture.* We replaced the standard classification head with a custom 1D Convolutional
 308 Neural Network (CNN) to capture local dependencies in code. The architecture consists of:

- 309 (1) **Encoder:** CodeBERT outputs a sequence of vectors ($L \times 768$).
- 310 (2) **Convolution:** A 1D Conv layer (256 filters, kernel size 3) scans the sequence for local risk
 311 patterns.
- 312 (3) **Activation:** LeakyReLU ($\alpha = 0.1$) is applied to maintain gradient flow.
- 313 (4) **Pooling:** Global Max Pooling extracts the most significant features.
- 314 (5) **Classification:** A dense layer maps the pooled features to the binary output (Backport/Reject).

316 *E.3.3 Training Dynamics.* The model was trained for 100 epochs using the AdamW optimizer
 317 ($lr = 2e^{-5}$). We observed peak generalization performance between Epochs 14 and 40 ($F1 \approx 0.80$).
 318 Beyond this point, the model exhibited clear signs of overfitting, where training loss converged
 319 to near-zero (< 0.01) while validation accuracy declined. The best-performing checkpoint was
 320 selected for the final evaluation.

321 E.4 Dual-Stream Architecture: RoBERTa + CodeBERT

323 A significant limitation of standard Transformer models is the 512-token context window. In complex
 324 backports, the combination of a detailed commit message and the associated code diff often exceeds
 325 this limit, leading to truncation and information loss. To mitigate this, we implemented a **Dual-
 326 Stream Multi-Modal Architecture** that processes natural language (intent) and programming
 327 language (implementation) in parallel.

328 *E.4.1 Architecture Design.* The system consists of two independent encoders that do not share
 329 weights:

- 330 (1) **Text Stream (Intent):** The commit subject and message are tokenized (max length 128)
 331 and passed through `roberta-base`. We utilize the [CLS] token embedding to capture the
 332 semantic intent of the change (e.g., distinguishing a "Bug Fix" from a "Feature Request").
- 333 (2) **Code Stream (Implementation):** The smart-compressed diff is tokenized (max length 400)
 334 and passed through `microsoft/codebert-base`. This stream focuses on identifying risky
 335 technical patterns (e.g., dependency changes or high-churn modifications) independent of
 336 the author's description.

338 *E.4.2 Feature Fusion.* The output vectors from both streams, $v_{text} \in \mathbb{R}^{768}$ and $v_{code} \in \mathbb{R}^{768}$, are
 339 concatenated to form a unified representation $v_{fused} \in \mathbb{R}^{1536}$. This vector is passed through a
 340 Multi-Layer Perceptron (MLP) with dropout ($p = 0.3$) for final classification.

$$341 \quad 342 \quad y = \sigma(W_2 \cdot \text{ReLU}(W_1 \cdot [v_{text} \oplus v_{code}] + b_1) + b_2) \quad (2)$$

344 E.4.3 *Results.* This architecture achieved an accuracy of **76.6%**, significantly outperforming the
345 single-stream CodeBERT baseline (58% on raw data). This confirms that decoupling the represen-
346 tation of *what was said* (RoBERTa) from *what was done* (CodeBERT) allows the model to retain
347 critical context that would otherwise be lost to truncation.

348
349 **References**

- 350 [1] Rakesh Agrawal and Ramakrishnan Srikant. 1994. Fast algorithms for mining association rules. In *Proceedings of the*
351 *20th International Conference on Very Large Data Bases (VLDB)*. Morgan Kaufmann, Santiago, Chile, 487–499.
352 [2] Sergey Brin, Rajeev Motwani, and Craig Silverstein. 1997. Beyond market baskets: Generalizing association rules to
353 correlations. In *Proceedings of the 1997 ACM SIGMOD International Conference on Management of Data*. ACM, New York,
354 NY, USA, 265–276.
355 [3] Bernhard Ganter and Rudolf Wille. 1999. *Formal Concept Analysis: Mathematical Foundations*. Springer-Verlag, Berlin,
356 Heidelberg.

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384

385

386

387

388

389

390

391

392