

Supplementary Material for BackportCheck: An Automated Risk Reporting Tool for Code Review

A Feature Engineering Taxonomy

Table 1 details the comprehensive list of 37 features extracted by the inference engine. To ensure clarity, each feature is defined individually with its specific rationale for risk assessment.

B Explainability Prompt Structure

To ensure consistent and trustworthy explanations, the raw feature vector and XGBoost predictions are injected into the structured template below (Figure 1) before being processed by the LLM.

```

SYSTEM: Act as a Senior OpenStack Release Manager. Justify the decision to {VERDICT} this backport.
== DECISION ==
VERDICT: {VERDICT} (Confidence: {PROB}%, Threshold: {THRESHOLD})
CATEGORY: {UI_DISPLAY_TYPE}
== CONTEXT: HOW TO INTERPRET THE DATA ==
- Author Trust Score: Ratio of accepted backports. 0.0=New, >0.5=Trusted.
- Historical File Prob: Probability these specific files are usually backported.
- Change Entropy: Code complexity (0=Simple, >4=Complex/Scattered).
- Churn Density: Lines changed per file (High = Dense/Risky).
- Modifies DB/API: Critical risk factors.
== FULL FEATURE VECTOR (Internal Data) ==
{FEATURE_VECTOR_JSON}
== CHANGE ARTIFACTS ==
Commit Message: "{COMMIT_MESSAGE}"
Files Modified: {FILE_LIST_STRING}
== INSTRUCTIONS ==
Write a professional, 2-3 sentence justification.
(1) Analyze the Vector: Look at the "Full Feature Vector" above. Find the anomalies or strong signals.
(2) Synthesize: Do not list the numbers. Explain their meaning.
(3) Explain the Verdict:
    • If REJECTED: Is it the Author? The Complexity? The specific Files?
    • If ACCEPTED: Is it the Safety (Config/Doc)? The High Trust?
RESPONSE:

```

Fig. 1. **The Prediction-Aligned Prompt Template.** The system dynamically populates the placeholders (in brackets) with the inference results and the raw feature vector before sending the request to the LLM.

C Validation via Formal Concept Analysis (FCA)

To validate the utility of our features and justify the dashboard design, we applied Association Rule Mining (Apriori algorithm) on the historical dataset of 3,422 changes. Table 2 presents the statistically dominant decision rules (Confidence = 1.0, Lift > 1) that emerged from the data. These rules were grouped into "Concept Families" to structure the BackportCheck user interface.

D Interface Visualization

Figure 2 illustrates the client-side integration of BackportCheck within the Gerrit environment. The interface is structured into five functional zones to facilitate transparent decision-making:

The workflow begins with the **Decision Banner (A)**, which provides an immediate probabilistic verdict based on the model's score. This signal is contextualized by the **AI Explanation (B)**, where the language model synthesizes the risk factors into a natural language justification. For verification, the **Risk Dashboard (C)** exposes the underlying quantitative metrics, such as Author Reliability, File Frequency, and Code Spread. Additionally, the interface includes **Sensitivity Controls (D)**, allowing users to calibrate the model's strictness, and **Metric Definitions (E)** to ensure the calculation logic remains transparent to the maintainer.

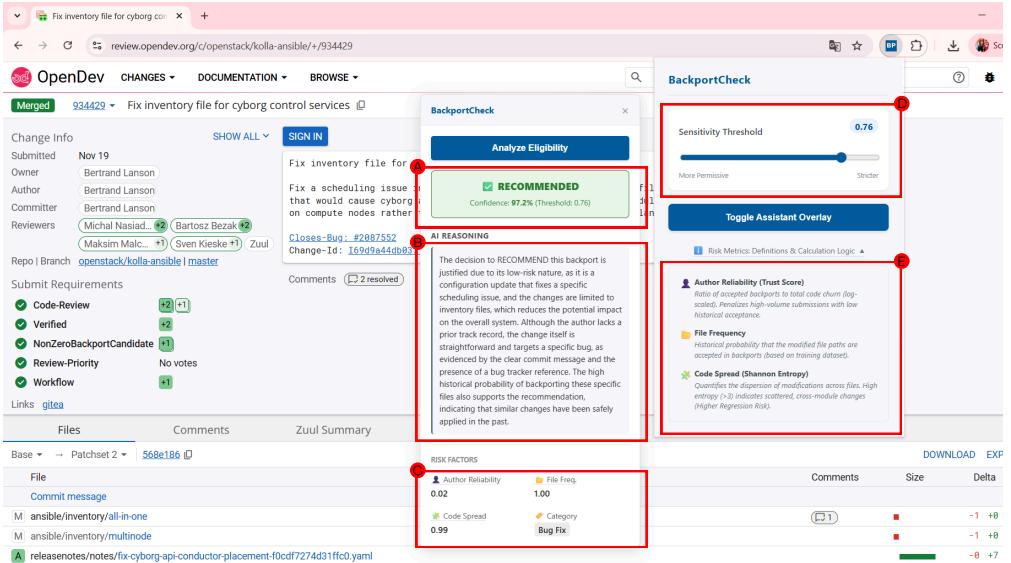


Fig. 2. **Annotated BackportCheck Interface.** The overlay augments the Gerrit code review screen with five decision-support layers: (A) The probabilistic verdict banner, (B) Natural language AI reasoning, (C) The quantitative risk dashboard, (D) Sensitivity threshold controls, and (E) metric definitions for transparency.

Table 1. Comprehensive Taxonomy of the 37 State-Aware Features. The feature vector is constructed to capture both the technical complexity of the code and the socio-technical context of the submission.

Dimension	Feature Variable	Definition & Rationale
1. Historical Context (Socio-Technical State)		
Author Reputation	author_trust_score	Calculated as $T_a = \frac{N_{accepted}}{\log(1+\sum C_{churn})}$. Penalizes massive code dumps while rewarding consistent accepted patches.
	author_success_rate	The ratio of the author's previously accepted backports to total submissions.
	author_submission_count	Indicates alignment with policy.
Environmental Bias	project_acceptance_rate	The total number of past changes submitted. Serves as a proxy for the author's experience level.
	historical_file_prob	The baseline acceptance rate of the specific sub-project (e.g., Nova vs. Horizon), used to normalize performance. Bayesian probability estimate: "How often have changes to this specific file path been backported in the past?"
2. Complexity & Volume (Entropy)		
Scatter Metrics	change_entropy	Shannon Entropy ($H(X)$) of change dispersion. High entropy (> 5) implies "Spaghetti Code" (scattered logic).
	safe_entropy_interaction	Interaction term correlating low entropy with safe file types (e.g., config files).
Cognitive Load	churn_density	Average lines changed per file. High density implies complex logic concentrated in single modules.
	churn_log_size	Logarithmic scale of total lines changed, smoothing the impact of massive refactorings.
	file_count	Total number of modified files. A proxy for the "blast radius" of the change.
	deletion_ratio	Ratio of deleted lines to total churn. High values indicate cleanup/dead-code removal (lower risk).
3. Semantic Intent (NLP & Regex)		
Classification	is_fix	Boolean: Commit message explicitly references a bug fix pattern.
	is_feature	Boolean: Commit introduces a new feature (typically rejected in stable backports).
	is_refactor	Boolean: Commit indicates a refactoring operation without functional changes.
	is_revert	Boolean: Commit is a revert of a previous change (high priority, generally low risk).
	is_maintenance	Boolean: Routine maintenance tasks like version bumps or dependency upgrades.
	is_deployment	Boolean: Changes to deployment-specific scripts/playbooks.
	msg_readability_ease	Flesch Reading Ease score. Higher scores indicate clearer, simpler justifications.
	msg_gunning_fog	Gunning Fog index estimating years of education needed to understand the text.
	references_bug_tracker	Boolean: Message links to an external issue tracker (e.g., Launchpad/Jira).
Clarity	has_subject_tag	Boolean: Compliance with project subject-tagging guidelines (e.g., [Nova]).
4. Domain Risk & Safety (Heuristics)		
High Risk	modifies_db_migration	Boolean: Changes to database schema migration files (Alembic). Critical regression risk.
	modifies_public_api	Boolean: Changes to public API definitions. High risk of breaking backward compatibility.
	modifies_dependencies	Boolean: Changes to dependency files ('requirements.txt'). Risk of downstream breakage.
Value Driver	has_security_impact	Boolean: Presence of security keywords (CVE, vulnerability). High value overrides risk.
Safety	is_pure_config	Boolean: Change is exclusively limited to configuration files ('.conf', '.ini').
	is_doc_only	Boolean: Change modifies only documentation files ('.rst', '.md').
	modifies_config	Boolean: Configuration files are involved (even if mixed with code).
5. Structural Topology & Meta Context		
Architecture	directory_depth	Maximum directory depth of modified files. Deep nesting often implies core logic.
	file_extension_entropy	Heterogeneity of file types (e.g., mixing SQL, Python, and Bash in one commit).
Composition	config_line_ratio	Proportion of lines changed that belong to configuration files.
	code_line_ratio	Proportion of lines changed that belong to source code (Python/C++).
	is_test_change	Boolean: Change is confined to test suites.
	is_ci_change	Boolean: Changes to Continuous Integration pipelines (Zuul/Tox).
Metadata	is_bot	Boolean: Commit generated by automated bots (e.g., translation syncs).
	has_gerrit_topic	Boolean: Change is part of a named topic series in Gerrit.
	is_deployment_project	Boolean: Repository belongs to a deployment tool (e.g., Kolla, Ansible).

Table 2. Key Association Rules Derived from Historical Data. These stable patterns justify the grouping of features into the "Risk" and "Safety" dashboard panels.

Antecedents (Contextual Signals)	Outcome	Conf.	Lift
<i>Risk Patterns (Rejection Signals)</i>			
Low Success Rate + Low Readability + Is BugFix + High File Count	Rejected	1.00	31.4
Low Code Ratio + Med Trust + High File Count + Is BugFix	Rejected	1.00	30.8
High Entropy + Deep Directory + New Author	Rejected	1.00	28.5
CI Change + Low Readability + Modifies DB	Rejected	1.00	25.1
<i>Intrinsic Safety (Configuration & Documentation)</i>			
Config Only + Low Entropy + High Trust + Med Readability	Accepted	1.00	5.9
Low File Count + Config Only + High Success Rate	Accepted	1.00	5.9
Low Code Ratio + Config Only + High Trust	Accepted	1.00	5.9
<i>Logic Safety (Trusted Code Patterns)</i>			
High Trust + Med Readability + Low Entropy + Low File Count	Accepted	1.00	5.9
High Success Rate + Low Directory Depth + Med Readability	Accepted	1.00	5.9
High Trust + Med Config Ratio + Low Entropy	Accepted	1.00	5.9